

Supplementary Information for “Certified cost forecasts and refutable regime maps for exact quantum compilation”

Sze Chun Yiu

Stockholm University

1 Formal model definitions

1.1 Pauli notation

An n -qubit Pauli word modulo global phase is represented by $(x, z) \in \mathbb{F}_2^{2n}$. Multiplication is vector addition. Its support is $\text{wt}(x, z) = |\{q : (x_q, z_q) \neq (0, 0)\}|$, and its commutation parity with (x', z') is $\langle (x, z), (x', z') \rangle = x \cdot z' + z \cdot x' \pmod{2}$. All theorem statements below quantify over every n and every target tuple for which the stated feasible set is nonempty.

1.2 Shared-tag paired-block model

An instance is six nonidentity targets P_{jk} , with $j \in \{0, 1, 2\}$ and $k \in \{0, 1\}$. A configuration consists of six nonidentity frames R_{jk} , a tag S , and a central index $c_j \in \{0, 1\}$ for each block, subject to

$$\langle R_{j0}, R_{j1} \rangle = 1, \quad (1)$$

$$\langle S, R_{j0} \rangle = \ell_0, \quad \langle S, R_{j1} \rangle = \ell_1, \quad (2)$$

$$(\ell_0, \ell_1) \in \{(0, 1), (1, 0)\}, \quad (3)$$

with the same (ℓ_0, ℓ_1) for all three blocks. Write $T_{jk} = P_{jk}R_{jk}$. For local Pauli letters define

$$F_3(a, b, c) = \begin{cases} 1, & a = b = c \neq I, \\ \mathbf{1}[a \neq I] + \mathbf{1}[b \neq I] + \mathbf{1}[c \neq I], & \text{otherwise.} \end{cases} \quad (4)$$

The structural objective is

$$\begin{aligned} C_{\text{shared}} = & \sum_{j=0}^2 \left[2(\text{wt}(R_{j,c_j}) - 1) + 4(\text{wt}(R_{j,1-c_j}) - 1) \right] + 2 \text{wt}(S) \\ & + \sum_{q=1}^n \sum_{k=0}^1 F_3(T_{0k}[q], T_{1k}[q], T_{2k}[q]). \end{aligned} \quad (5)$$

The subtracted one in each frame term is a constant normalization over feasible configurations. The support coordinate is $\max_{j,k} \text{wt}(R_{jk})$.

Sze Chun Yiu: sze-chun.yiu@fysik.su.se

1.3 Rank-two dependent-triple model

An instance contains two ordered target triples P_{bk} , $b \in \{A, B\}$ and $k \in \{0, 1, 2\}$. A configuration selects a relative permutation π_B , a central branch $c_b \in \{0, 1, 2\}$ in each block, independent frames R_{b0}, R_{b1} satisfying $\langle R_{b0}, R_{b1} \rangle = 1$, dependent frames $R_{b2} = R_{b0}R_{b1}$, and two shared tags S_0, S_1 . Define

$$\lambda(R) = 2\langle S_0, R \rangle + \langle S_1, R \rangle \in \{0, 1, 2, 3\}. \quad (6)$$

Feasibility requires

$$(\lambda(R_{A0}), \lambda(R_{A1})) = (\lambda(R_{B0}), \lambda(R_{B1})) = (a, b), \quad a, b \in \{1, 2, 3\}, \quad a \neq b. \quad (7)$$

The dependent label is then $a \text{ xor } b$, the third nonzero label. With π_A the identity, the objective is

$$C_{\text{rank2}} = \sum_{b \in \{A, B\}} \left[t_c (\text{wt}(R_{b, c_b}) - 1) + t_{nc} \sum_{k \neq c_b} (\text{wt}(R_{bk}) - 1) + t_r \sum_{k=0}^2 \text{wt}(P_{b, \pi_b(k)} R_{bk}) \right] + t_{\text{tag}} (\text{wt}(S_0) + \text{wt}(S_1)). \quad (8)$$

The unit structural objective is $(t_c, t_{nc}, t_{\text{tag}}, t_r) = (2, 4, 2, 1)$. The support coordinate is $\max_{b, k \in \{0, 1\}} \text{wt}(R_{bk})$.

1.4 Six-term linear-combination model

An instance is a formal batch $(p_1, \dots, p_6)$ of nonidentity n -qubit Pauli words; repeated words are allowed. A configuration is a set partition G of $\{1, \dots, 6\}$, a factoring bit φ_g for every block, and either shared or dedicated index ancillas. For a block g of size m , let $b(m) = \lceil \log_2 m \rceil$, with $b(1) = 0$. Its common factor F_g contains the local letter shared nontrivially by every p_i , $i \in g$, and identity elsewhere; $r_i = p_i F_g$ modulo phase. Let $f = 1$ when G has at least two blocks and $f = 0$ otherwise. The selection contribution of g is

$$\text{SEL}_g = \begin{cases} \sum_{i \in g} (f + b(m) + 1) \text{wt}(p_i), & \varphi_g = 0, \\ (f + 1) \text{wt}(F_g) + \sum_{i \in g} (f + b(m) + 1) \text{wt}(r_i), & \varphi_g = 1. \end{cases} \quad (9)$$

Writing $d(1), \dots, d(6) = (0, 0, 1, 2, 4, 6)$ for the frozen balanced-tree depth sums,

$$\text{PREP}(G) = \mathbf{1}[|G| \geq 2] (2|G| - 3) + \sum_{g: |g| \geq 2} [(|g| - 1)(1 + f) + d(|g|)], \quad (10)$$

$$\text{WIDTH}(G) = \mathbf{1}[|G| \geq 2] |G| + \begin{cases} \max_g b(|g|), & \text{shared,} \\ \sum_g b(|g|), & \text{dedicated.} \end{cases} \quad (11)$$

The exact family cost is the minimum of $\sum_g \text{SEL}_g + \text{PREP} + \text{WIDTH}$ over all 203 partitions, factoring bits and ancilla choices. The reference is the six-singleton partition, with cost $C_U = 2 \sum_i \text{wt}(p_i) + 15$.

2 Forecast families and explanation closure

For the shared-tag model, let C_{ref} denote the frozen reference construction, $C_{\leq 1}$ the exact minimum over the support-one restricted family, $C_{\leq 2}$ the exact minimum over all support-at-most-two frames, and C_{DP} the unrestricted dynamic-programming minimum. The

archived implementation names these quantities C_{R6L} , C_{D+} , C_{D++} and C_{DP} , respectively. The cost theorem proved in Section 3 is

$$C_{DP} = C_{\leq 2} \quad (12)$$

for every admitted size under Eq. (5).

The explanation families are nested by mechanism rather than by evidence status. B contains the initial central-borrow construction. B' additionally permits the borrow home to lie outside its block target support or inside the union support. B'' contains the verified split-and-borrow hybrid shapes. Their cost functions are the exact minima within the declared families, not fitted regressors.

The first explanation attack evaluates all 740 frozen instances, not only strict-gap rows. Relative to $f_{B'}$, exact offsets have counts 389, 325, 13, 13 at offsets 0, 1, 2, 3. Six of seven $f_{B'}$ fibres are cost-mixed. Selecting the 64 rows with $C_{\leq 2} < \min(C_{\leq 1}, f_{B'})$ instead yields a uniform one-unit gap and no mixed fibre, proving that the apparent regularity is induced by selection.

The last commuting-support-two obligation is discharged by an induction, and because that induction is the composition step of the all-size envelope, we state it in full. The envelope is proved by a twelve-link reduction chain committed in the QG-7e receipt. Each link replaces a class of configurations by an explicit alternative while preserving three invariants: feasibility, meaning every anticommutation and tag-syndrome constraint of the model; non-increase of the objective, $\Delta \leq 0$ for every produced alternative; and size independence, because every edit acts at named local coordinates only, so a domination fact proved once holds at every admitted size. The links ahead of the geometry step are complete-domain checks, with lemma names following the committed receipts: the all- n exchange theorem $C_{DP} = C_{\leq 2}$ of Section 3; the MG mirror, an exact cost-preserving involution of the whole configuration, verified on its complete 2^{24} -state configuration domain; per-qubit letter permutations as cost gauges, verified on the complete 6×2^{24} domain; the receipt-bound eliminations of support-at-least-three frames, (2, 2) blocks, class-(0, 0) qubits and out-of-support tag letters on complete domains of 768, 27,216, 8, 18,432, 43,688 and 1,440 cases; the M1 inventory on its complete 262,144-case domain, which classifies every irreducible block as anchored, phantom or commuting-support-two; the T1 tag-letter prune with exact refund two over 12,288 cases; the T4b pinned census over its complete 2^{29} -state domain, which leaves exactly 135,604 residual patterns; and a complete 5,342,016-row re-derivation, with zero mismatches over $n \in \{1, 2, 3\}$, of the per-block target permutation as a degree of freedom of the committed enumerator.

The base case of the induction is the per-geometry domination lemma. The 135,604 census patterns are re-derived verbatim and dispatched, with zero patterns left open, into 378 geometries, which are exactly the unordered pairs, with repetition, of the 27 registered irreducibility roles: $27 \cdot 28/2 = 378$. For each geometry, the complete raw target-letter domain at the two active commuting-support-two qubits has 2^{24} assignments, and every one of them admits an alternative with cost change $\Delta \leq 0$ and strictly fewer commuting-support-two blocks, scored over the MG mirror and all $2^3 = 8$ per-block target assignments obtained by permuting the two targets independently within each remaining block. The committed per-geometry table records 378 closed geometries and residue zero, so the base case covers all $378 \times 2^{24} = 6,341,787,648$ states without exception.

The composition step is an induction on the number of commuting-support-two blocks, which serves as its strictly decreasing measure: one application of the geometry lemma removes exactly one commuting-support-two block and never creates one, so the iteration terminates at a commuting-support-two-free irreducible configuration. A fact proved

on the fixed per-geometry letter domain transfers to composed configurations at every size because each scored alternative changes letters only at the two commuting-support-two qubits and reads the state only through a per-block permutation of a block's two targets, which moves no support; the cost difference computed inside the lemma is therefore the exact global cost difference in every context in which the geometry occurs. Cost never increases along the iteration, so the image of an unrestricted optimum is a configuration of no larger cost inside a terminal family: the T3 consolidation on its complete 14,680,064-case domain moves weight-three-tag commuting-support-two-free configurations into B' and B'' shapes at $\Delta \leq 0$, and T5 together with the M1 structure maps the terminal shapes into the committed grammars underlying $C_{\leq 1}$, $f_{B'}$ and $f_{B''}$. This proves $\min\{C_{\leq 1}, f_{B'}, f_{B''}\} \leq C_{\text{DP}}$. The reverse inequality is definitional, because each named family is a subfamily of the support-at-most-two family, so $\min\{C_{\leq 1}, f_{B'}, f_{B''}\} \geq C_{\leq 2} = C_{\text{DP}}$. Hence

$$C_{\text{DP}} = \min\{C_{\leq 1}, f_{B'}, f_{B''}\}. \quad (13)$$

Two alternative routes to this statement are dead, so the per-block-permutation induction is the only registered discharge. The single-pinner normalization is accepted all- n on its local product domain of 133,349,376 states, with deterministic replay byte-identical, but its registered composition fields remain `CHAIN-ALL-N = false` and `GLOBAL-B''-COMPLETENESS = false`: the link from the local normalization to a global chain is unproved. The two-coordinate chain representation is refuted rather than open: an exact three-Tag configuration of tag weight three in which two of the three blocks are irreducible commuting-support-two, with no reducible class-(0,0) coordinate in either, falsifies its premise. The successful proof uses per-block target permutations and converts neither negative into support. Machine-checkability of the composition step is bounded by what the per-geometry audits record: the committed receipts state the closure row by row (378 closed geometries, residue zero, 135,604 of 135,604 census patterns dispatched, and the dead-route flags above), a receipt-level checker in the artifact re-derives this arithmetic, and the underlying 6,341,787,648-state enumeration is the scientific run itself, not re-executed by the checker.

3 All-size normalization proofs

3.1 Shared-tag support-two proof

Fix a feasible configuration and a frame $R = R_{jk}$ with support $w \geq 3$; let R' be its anticommuting partner. At each $q \in \text{supp}(R)$ assign

$$(\alpha_q, \beta_q) = (\langle R[q], R'[q] \rangle, \langle S[q], R[q] \rangle) \in \mathbb{F}_2^2. \quad (14)$$

The total α parity is one. Any such multiset of size at least three contains either a (0,0) singleton or an equal pair. Indeed, if there is no zero and no repetition, it is a subset of the three nonzero vectors; the only size-three case has zero total α , a contradiction. Removing the indicated proper subset Q , with $1 \leq |Q| \leq 2$, therefore changes neither total α nor total β . Zeroing R on Q preserves block anticommutation and the shared tag syndrome, leaves the other two blocks unchanged, and needs no tag repair.

It remains to compare cost. For each removed local letter, central or noncentral frame support refunds respectively two or four units. Exhaustive evaluation of the local factor F_3 in Eq. (5) covers the removed nonidentity letter, its partner and tag letters, the raw target letter, the two other block letters, the branch position and both multipliers. Across all

18,432 cases the increase in F_3 is at most the frame refund. Thus the edit is non-increasing in cost and strictly decreases total frame support. Repeating it terminates with support at most two. The four two-letter class patterns that do not contain a removable zero-sum proper subset explain why this argument stops at two rather than one.

3.2 Rank-two support-one proof

The support-one proof begins from a support-two normal form for the same feasible family. The precursor certificate uses a syndrome-preserving relabel-and-delete grammar. Its full local action space contains 612 support-three profile cases; 21 are unsafe before imposing Eq. (7), and none of those 21 is feasible after the complete label predicate is applied. Hence every optimum has a support-two representative. The packaged checker reconstructs these profiles from local Pauli descriptors and verifies the acceptance filter; the 144-case support-two control retains 36 feasible obstructions to that older edit grammar, so the precursor alone does not claim support one.

Now take any support-two representative. Each anticommuting pair has exactly one qubit at which its two local letters anticommute; call it the core. Delete every non-core active column, align each remaining local basis at its core, and reconstruct the tags only after both blocks have been localized. The complete local bounds are:

1. deleting a commuting active column refunds at least four objective units (and deleting an anticommuting column refunds at least seven);
2. changing an anticommuting local basis at a fixed core changes the restoration term by at most three while leaving the frame contribution equal to ten;
3. on a shared core, feasibility forces the two blocks to use the same local basis, and the canonical dual tag costs four;
4. on distinct cores, the canonical dual tags cost eight; if both blocks were already support one the original tag costs at least eight, while otherwise one non-core deletion from each affected side supplies the additional cost.

The numerical inequalities $4 > 3$, $4 + 4 \geq 8$ and the tag lower bounds close the same-core and distinct-core cases. The edit preserves Eq. (7) by construction and never increases Eq. (8). Both independent generators end on one core. Support zero is infeasible because $\langle I, I \rangle = 0$, so the anticommutation constraint cannot be met. Hence the intrinsic support number is one.

Table 1: Finite obligations supporting the two all-size normalization results. These enumerations discharge local cases inside the analytic compositions; they are not samples from a deployment population.

Model	Obligation	Rows	Outcome
Shared-tag Pauli	local exchange inequality	18,432	0 violations
Shared-tag Pauli	parity-class tuples	43,688	exact support-two boundary; 0 failures for support 3–8
Rank-two dependent triple	precursor action profiles	612	21 unsafe before full acceptance; 0 accepted unsafe
Rank-two dependent triple	deletion cases	2,880	worst changes -4 and -7
Rank-two dependent triple	core alignment	6,912	maximum increase $+3$
Rank-two dependent triple	same-core tag rigidity	576	0 different-basis counterexamples
Rank-two dependent triple	distinct-core tag lower bound	9,216	minimum tag cost 8

4 Objective-indexed certificates

4.1 Shared-tag support-two cone

In the shared-tag descent, deleting a central or noncentral frame letter refunds respectively t_c or t_{nc} , while the local restoration functional can increase by at most $2t_r$. The same exchange therefore proves an all-size support-two normal form throughout

$$t_c \geq 2t_r, \quad t_{nc} \geq 2t_r. \quad (15)$$

The unit point $(t_c, t_{nc}, t_{\text{tag}}, t_r) = (2, 4, 2, 1)$ lies on the central boundary. The registered outside point $(1, 7, 4, 3)$ contains exact support-three optima with $C_{\text{DP}} = 11 < C_{\leq 2} = 13$. It refutes an objective-independent support-two theorem. Because violation of a sufficient inequality is not a converse proof, the global sharpness of the cone remains open.

4.2 Rank-two support-one region

For a deleted commuting column, let (r_c, r_{nc}, Δ_r) record refunded central support, refunded noncentral support and restoration increase. The complete 1,728-row resource domain has two undominated worst vectors, $(0, 2, 2)$ and $(1, 1, 2)$. Core alignment may add a further $3t_r$, while moving the two tag words to a second core may add $2t_{\text{tag}}$. Requiring each worst refund to cover either obligation gives

$$2t_{nc} \geq 5t_r, \quad t_c + t_{nc} \geq 5t_r, \quad (16)$$

$$2t_{nc} \geq 2t_r + 2t_{\text{tag}}, \quad t_c + t_{nc} \geq 2t_r + 2t_{\text{tag}}. \quad (17)$$

For $(2, 4, 2, 1)$ the left-minus-right margins in displayed order are $(3, 1, 2, 0)$. For $(3, 5, 2, 1)$ they are $(5, 3, 4, 2)$. Thus the first point lies on the certificate boundary and the second lies strictly inside. These inequalities are sufficient for this proof. Their violation provides no converse statement.

5 Six-term pair-gain equivalence

For $B \subseteq \{1, \dots, 6\}$, write $s(B) = \sum_{i \in B} \text{wt}(p_i)$ and let $f(B)$ count all-equal nonidentity columns. The gain of a factored block of size m has coefficients (A_m, B_m) multiplying $(f(B), s(B))$:

$$(A_m, B_m) = (0, 0), (4, -1), (10, -2), (14, -2), (23, -3) \quad (m = 1, \dots, 5). \quad (18)$$

For $m = 6$ the gain is $23f(B) - 2s(B) + 1$. A partition contributes the sum of its block gains plus the shape constant shown in Table 2.

Let $g(e) = 4f(e) - s(e)$ for each pair. Condition P_0 consists of $g(e) \leq 0$ for every pair, $g(e) + g(f) + 1 \leq 0$ for every two disjoint pairs, and $g(e) + g(f) + g(h) + 2 \leq 0$ for every perfect matching. Summing pair clauses within a block yields $s(B) \geq 6f(B)$ for triples, $s(B) \geq 8f(B)$ for quads, $s(B) \geq 10f(B)$ for quints, and $s(B) \geq 12f(B) + 2$ for the six-term block. Substitution makes every non-pair shape gain nonpositive. The three multi-pair shapes are exactly the three clauses of P_0 . Conversely, violation of any clause makes its corresponding pair partition strictly cheaper than the singleton reference. This proves $C_F = C_U$ if and only if P_0 for every admitted batch and every n .

Table 2: Complete shape ledger for the six-term equivalence. T_m is the affine gain of one factored block of size m ; g_i denotes a pair gain.

Shape	Partitions	Exact gain form	Bound under P_0
$1 + 1 + 1 + 1 + 1 + 1$	1	0	equality
$2 + 1 + 1 + 1 + 1$	15	g_1	pair clause
$2 + 2 + 1 + 1$	45	$g_1 + g_2 + 1$	two-disjoint-pair clause
$2 + 2 + 2$	15	$g_1 + g_2 + g_3 + 2$	perfect-matching clause
$3 + 1 + 1 + 1$	20	$T_3 - 1$	$T_3 \leq 0$
$3 + 2 + 1$	60	$T_3 + g_1$	both nonpositive
$3 + 3$	10	$T_3 + T'_3$	both nonpositive
$4 + 1 + 1$	15	$T_4 - 1$	$T_4 \leq 0$
$4 + 2$	15	$T_4 + g_1$	both nonpositive
$5 + 1$	6	$T_5 - 3$	$T_5 \leq 0$
6	1	$23f - 2s + 1$	at most $-f - 3$

6 State-preparation model and exact feature construction

A state is the full signed stabilizer group in canonical sorted binary form. The allowed gates are H , S , $S^\dagger$ and controlled-NOT with costs 1, 1, 1, 3. The exact referee is Dijkstra search from $|0 \dots 0\rangle$ over the complete stabilizer-state graph. The deterministic reference synthesis processes qubits in ascending order, selects the lexicographically smallest admissible X pivot, uses the X elimination route when available, otherwise uses the forced signed- Z route, and applies the inverse elimination sequence for preparation. Sign repairs and controlled-NOT routes are therefore deterministic functions of the target state.

The 127 coordinates are ordered in three blocks. The 33 schedule/tensor coordinates comprise 13 base quantities and 20 trace quantities. The base quantities are five reference event counts, reference cost, binary X rank, tensor-factor count, the lower bound $r_X + 3(n - c)$, and four declared differences from these quantities. The trace block records maximum, first and last step costs, the maximizing index, cost descents and moment, high/zero-cost step counts, event maxima and sign/ Y co-occurrences, signed-pivot and pivot-weight summaries, Z -route and rank-drop counts, largest tensor-factor size and reference cost, the enlarged-schedule cost, and its improvement over the reference.

The 53 path coordinates comprise four gate totals; minimum, maximum, sum of squares and zero count for each of five per-qubit gate channels; eight directed controlled-NOT graph summaries; all 16 ordered adjacent gate-kind transition counts; and first kind, last kind, run count, maximum run and number of distinct kinds. The 41 state coordinates comprise the negative-sign count; positive/negative Pauli-weight histograms for weights zero through four; positive/negative Y -count histograms; positive/negative binary parity-class counts; and minimum, maximum, sum of squares and zero count of the per-qubit X , Y and Z column marginals. The archive gives the exact coordinate names, order and formulas and supplies every panel vector.

The complete $n \leq 3$ domain has 1,146 states and 1,109 feature cells, including 1,072 singletons. No cell is mixed, so a lookup recovers every in-domain label. The ratio $1,109/1,146 = 0.967714$ shows that the representation gives little compression.

Table 3: Complete and prospective state-preparation counts retained in the analysis.

Domain	States	Reference compiler exact
$n = 1$ complete	6	5
$n = 2$ complete	60	28
$n = 3$ complete	1,080	189
$n = 4$ frozen panel	120	32

7 Frozen panel, transfer rule and shuffle null

Using NumPy’s default generator with seed 20260821, each candidate panel state is produced from $|0000\rangle$ by 24 draws. A draw chooses uniformly among H , S , $S^\dagger$, and controlled-NOT; a one-qubit gate receives a uniform qubit, while a controlled-NOT receives a uniform control and then a uniform distinct target. Duplicate final states are discarded until 120 unique states have been retained in generation order. The review materials provide the ordered list of canonical states. Exact four-qubit costs are then read from the 36,720-state Dijkstra graph.

The representation and in-domain success criterion were frozen before four-qubit evaluation. Panel index parity defines two folds. For each test state, the other fold supplies counts for an exactly equal 127-coordinate vector. A seen cell predicts the strict majority label, with a tie mapped to negative; an unseen cell also predicts negative. Coverage is the number of test states whose vector occurs in the other fold. The 120 panel states form 119 cells: 118 singletons and one pure doubleton, so the partition invariant is $118 \times 1 + 1 \times 2 = 120$. Only the two states in the doubleton are covered across parity folds. Both are classified correctly; all 32 errors occur among the 118 uncovered states and equal the panel’s 32 positive labels.

For the null, the feature vectors, state order, folds, lookup rule and negative default remain fixed. Python’s pseudorandom generator with seed 20260828 produces 200 independent permutations of the 120 observed labels. Each permuted label vector is evaluated by the identical two-fold procedure. If $E_0 = 32$ is the observed error count and E_b is permutation b ’s error count, the reported empirical probability is

$$p_{\leq} = \frac{1}{200} \sum_{b=1}^{200} \mathbf{1}[E_b \leq E_0]. \quad (19)$$

The null counts range from 30 to 33, have mean 32.41, and give $p_{\leq} = 0.51$. This descriptive randomization result says that the declared lookup shows no improvement under this null on this panel. It is not an equivalence test.

The earlier prospective forecast is a separate frozen analysis. Its 120 panel states, predicted labels and predicted costs were fixed before the four-qubit referee was invoked. It matched 100 regime labels and 67 exact costs. Later feature work does not revise those predictions.

8 Nearest-work subtraction and consequence

Table 4: Claim-by-claim nearest-work subtraction. The comparison identifies the residual mathematical object and its consequence; it is not a priority certificate.

Advance	Closest work	Difference in object, assumptions and guarantee	Consequence retained here
Typed exact-regime record	Instance-space analysis and algorithm selection [1, 2]	Prior work maps features to performance or choices; the typed record separately represents theorem, certificate, representation and transfer fields with explicit missing-value semantics.	Numerical fields cannot borrow authority across type, objective or domain.
Shared-tag support two	Pauli compilation and cluster diagonalization [3, 4]	Prior work supplies Pauli compilation constructions; the present object is the declared three-block shared-tag grammar and the guarantee is an all-size support-two normal form under Eq. (S4).	Search may be restricted to support two in this grammar; no hardware claim follows.
Rank-two intrinsic support one	Verified and exact circuit optimization [5, 6]	Prior work establishes correct rewrites or exact matching; the present guarantee is tight support one for the dependent-triple grammar, with an objective-indexed proof region.	A conserved proof ceiling of five is not intrinsic support.
Six-term pair-gain boundary	Unitary partitioning and Pauli optimization [7, 8]	Prior work motivates grouping and Pauli optimization; the present object fixes six terms, 203 encodings and one structural objective, and proves an if-and-only-if three-clause boundary.	Reference exactness is decidable from pair-derived gains in this family only.
Adverse feature transfer	Quantum compiler prediction and bounded optimal Clifford synthesis [9–11]	Prior work predicts compilation choices or computes bounded optima; the present study couples a complete finite feature partition to a frozen next-size exact referee and retains failure.	Finite determination does not imply next-size coverage or transfer.

The comparison supports a narrow consequence. Instance-space analysis and compiler prediction remain the closest frameworks for feature-conditioned performance, while verified and exact synthesis remain the closest frameworks for correctness and bounded optimality. The present exact results add model-specific normal forms and a pair-gain equivalence; the typed record prevents their theorem authority from being transferred to the finite feature study. It does not displace the cited frameworks or establish a universal compiler theory.

9 Retained adverse and conditional results

The first prospective size-transfer forecast was fixed before the exact four-qubit distances were opened. It matched 100 of 120 reference-exactness labels and 67 of 120 exact costs. Both axes failed their exactness criteria. These outcomes remain visible because later representation work does not retroactively change a prospective forecast.

A separate search tested whether support two became strictly beneficial outside the proved support-one region. After correcting the treatment of infeasible local cells, the complete frozen search still found zero strict witnesses among 211,248 candidates under four objectives. The result excludes those candidates only. It does not prove that the objective region is globally sharp.

The measured gap between syndrome rank and intrinsic support equals a local exchange margin on two models. The equality has no residual degrees of freedom with two points, and it fails for one model when the syndrome rank is recomputed under a different aligned rewrite. It is retained as a diagnostic observation, not a theorem.

Table 5: Adverse and conditional results that constrain interpretation.

Question	Exact result	Allowed interpretation
Prospective state-preparation forecast	100/120 regime labels and 67/120 exact costs matched	the frozen forecast failed; no success-rate generalization follows
Support increase outside the certificate region	211,248 candidates, 0 strict witnesses under four objectives	a frozen-domain negative only; the global boundary remains open
Rank versus intrinsic support	slack equals the measured exchange margin on two models	a two-point, rewrite-dependent observation rather than a law
Rich state representation	1,109 cells for 1,146 states; 1,072 singleton cells	finite-domain vocabulary existence, not compact structure or transfer

10 Reproduction order

The review materials support four layers of checking. First, inspect the exact result records. Second, run the separate verifier, which checks theorem fields, objective vectors, panel-partition arithmetic, table values and required claim wording without invoking a scientific calculation. Third, regenerate all tables and the per-state feature and prediction record. Fourth, where the declared Python dependencies are available, rerun the selected scientific analyses and compare their deterministic outputs. These operations detect inconsistent values and implementation drift. They are same-team checks, not external replication or independent proof review.

References

- [1] John R. Rice. The algorithm selection problem. In *Advances in Computers*, volume 15, pages 65–118. Elsevier, 1976. DOI: [10.1016/S0065-2458\(08\)60520-3](https://doi.org/10.1016/S0065-2458(08)60520-3).

- [2] Kate Smith-Miles and Mario Andrés Muñoz. Instance space analysis for algorithm testing: Methodology and software tools. *ACM Computing Surveys*, 55(12):1–31, 2023. DOI: [10.1145/3572895](https://doi.org/10.1145/3572895).
- [3] Filipa C. R. Peres and Ernesto F. Galvão. Quantum circuit compilation and hybrid computation using pauli-based computation. *Quantum*, 7:1126, 2023. DOI: [10.22331/q-2023-10-03-1126](https://doi.org/10.22331/q-2023-10-03-1126).
- [4] Ewout van den Berg and Kristan Temme. Circuit optimization of hamiltonian simulation by simultaneous diagonalization of pauli clusters. *Quantum*, 4:322, 2020. DOI: [10.22331/q-2020-09-12-322](https://doi.org/10.22331/q-2020-09-12-322).
- [5] Kesha Hietala, Robert Rand, Shih-Han Hung, Xiaodi Wu, and Michael Hicks. A verified optimizer for quantum circuits. *Proceedings of the ACM on Programming Languages*, 5(POPL):1–29, 2021. DOI: [10.1145/3434318](https://doi.org/10.1145/3434318).
- [6] Raban Iten, Romain Moyard, Tony Metger, David Sutter, and Stefan Woerner. Exact and practical pattern matching for quantum circuit optimization. *ACM Transactions on Quantum Computing*, 3(1):1–41, 2022. DOI: [10.1145/3498325](https://doi.org/10.1145/3498325).
- [7] Artur F. Izmaylov, Tzu-Ching Yen, Robert A. Lang, and Vladyslav Verteletskyi. Unitary partitioning approach to the measurement problem in the variational quantum eigensolver method. *Journal of Chemical Theory and Computation*, 16(1):190–195, 2020. DOI: [10.1021/acs.jctc.9b00791](https://doi.org/10.1021/acs.jctc.9b00791).
- [8] Jennifer Paykin, Albert T. Schmitz, Mohannad Ibrahim, Xin-Chuan Wu, and A. Y. Matsuura. PCOAST: A pauli-based quantum circuit optimization framework. In *2023 IEEE International Conference on Quantum Computing and Engineering*, pages 715–726, 2023. DOI: [10.1109/QCE57702.2023.00087](https://doi.org/10.1109/QCE57702.2023.00087).
- [9] Nils Quetschlich, Lukas Burgholzer, and Robert Wille. Predicting good quantum circuit compilation options. In *2023 IEEE International Conference on Quantum Software*, pages 43–53, 2023. DOI: [10.1109/QSW59989.2023.00015](https://doi.org/10.1109/QSW59989.2023.00015).
- [10] Nils Quetschlich, Lukas Burgholzer, and Robert Wille. MQT Predictor: Automatic device selection with device-specific circuit compilation for quantum computing. *ACM Transactions on Quantum Computing*, 6(1):1–26, 2025. DOI: [10.1145/3673241](https://doi.org/10.1145/3673241).
- [11] Sergey Bravyi, Joseph A. Latone, and Dmitri Maslov. 6-qubit optimal Clifford circuits. *npj Quantum Information*, 8:79, 2022. DOI: [10.1038/s41534-022-00583-7](https://doi.org/10.1038/s41534-022-00583-7).