

Representational Compatibility Is Not Scientific Identity: Coordinate-Governed Mapping of Source-Local Scientific Projections

Working framework draft

August 2026

Abstract

Cross-domain scientific synthesis is vulnerable to false integration: identical names may denote different referents, nominally similar constructs may rest on incompatible operationalizations, and apparently contradictory claims may differ only in context, modality or attribution. This paper treats integration as a decision that has to be licensed rather than as a similarity score to be thresholded. Each source is represented as a source-local scientific projection bound to its provenance, and a proposed cross-source mapping records referent identity, construct identity, measurement and operationalization, context and representation equivalence as separate coordinates. Compatible projections may be glued into one global statement only under explicit preservation conditions; where a load-bearing coordinate is incompatible or unresolved the mapping remains a typed obstruction or a plural view rather than being forced into one canonical summary. Local exact worlds falsify the mapping semantics. In a disjoint, prospectively execution-frozen 32-case confirmatory holdout assembled from pinned external human and expert resources, coordinate-governed mapping made zero false merges against 0.1875 for flat predicate canonicalization; the paired difference was -0.1875 (95% bootstrap CI $[-0.34375, -0.0625]$). Against an exact-coordinate conservative control, the false-split difference was 0.000 with interval $[0.000, 0.000]$, satisfying the rule declared before those outputs were seen. A separate post-saturation battery then granted its comparator a strong semantic product carrying structured construct, measurement, context, provenance and missingness information, and tested only whether representational alignment is sufficient authority for scientific identity. On 400 prospectively frozen exact heterogeneous mapping contracts, the claim-relative compatibility layer decided 400/400 exactly against 250/400 for that strong product, a paired difference of $+0.375$ (95% domain-stratified bootstrap CI $[0.3275, 0.4225]$), with no false merges and complete clean-merge coverage. An ideal information-equivalent typed product given the same coordinates and rules tied at 400/400, which rules out an inherent-expressivity reading. These results concern already-structured scientific projections and exact successor contracts only. Outcomes are three-valued throughout: an outcome that could not be determined is reported separately from an outcome determined to be false, and the two are never merged. Raw-text extraction, live ontology-engine generality, newly commissioned expert agreement, recoverability in generated portraits and downstream utility all remain undetermined in that third sense.

Keywords: scientific knowledge integration; provenance; schema alignment; measurement harmonization; scientific pluralism; knowledge representation.

1 Introduction

Scientific claims are distributed across disciplines, each with its own terminology, measurement practices, contexts, and representational choices. A global synthesis can therefore fail even when every local statement is individually well formed: identical surface terms can refer to different objects, different labels can denote compatible objects, and apparently contradictory claims can differ only in modality, attribution, or context. The central problem of this paper is not retrieval itself, but the authority to identify two already-structured scientific representations as one global statement.

1.1 The proposal in brief

Each source is represented as a *source-local scientific projection* bound to its provenance. A proposed cross-source mapping records scientific coordinates such as referent, construct, measurement, context, polarity, modality, and attribution separately rather than reducing them to one similarity score. When the coordinates required by the claim are compatible, the projections may be *glued* into a single global statement. When a load-bearing coordinate is incompatible or unresolved, the mapping remains a typed *obstruction* or a *plural view* rather than being forced into a canonical merge.

That rule is evaluated at the layer for which external evidence currently exists. An initial public-reference study contains 32 already-structured cases derived from pinned external sources. A second, disjoint 32-case holdout was selected and execution-frozen before its system outputs, including its exact gold digest, source revisions, evaluator code, bootstrap seed, margins, and pass rule. On that confirmatory holdout, coordinate-governed mapping produced zero false merges against 0.1875 for flat predicate canonicalization; the paired difference was -0.1875 with a 95% bootstrap interval of $[-0.34375, -0.0625]$. The false-split difference against an exact-coordinate conservative control was 0.000 with interval $[0.000, 0.000]$, satisfying the rule declared before those outputs were seen.

The evidence is deliberately narrower than the mechanism’s design ambition. It does not establish raw-text extraction superiority, expert construct validity across the eight case families the broader design describes, reader-level recoverability of a generated portrait, downstream scientific-answer improvement, or necessity of every semantic coordinate. Each of those remains undetermined rather than being inferred from the mapping result.

The surviving empirical claim is the narrow one: for the frozen public-reference mapping task, explicit scientific-coordinate conditions plus an obstruction state reduce false merges relative to flat predicate canonicalization while satisfying the false-split guard declared in advance. The initial 32 cases are not pooled into the confirmatory verdict, and coordinate ablations with no measured effect are reported as coverage limits rather than as proof that those coordinates are unnecessary.

1.2 Structure of the paper

Section 2 positions the mapping rule among the extraction, schema-alignment, integration, stance and pluralism literatures it builds on. Section 3 defines source projections, mappings, gluing, obstruction, and the three-valued reporting discipline. Section 4 describes the two public-reference mapping datasets and their provenance. Section 5 records the prospective confirmatory design and its controls. Section 7 reports the frozen result and the ablations. Section 8 separates the supported mapping claim from what remains unmeasured, Section 9 states the bounded conclusion, and Section 10 lists every artifact behind a reported number.

2 Related work

The question asked here sits downstream of most of the machinery that cross-domain scientific integration usually needs. Extraction, schema matching, entity resolution, fusion, contradiction and stance analysis, and literature-based discovery are all treated as available capability. What remains is smaller and, we argue, unresolved by any of them: given two already-structured source projections, which coordinate differences license a merge, and which oblige the system to keep a provenance-bound obstruction visible instead.

2.1 Structured scientific knowledge bases and extraction

MUSE [6] constructs a full-text, multi-domain resource of source-grounded problem–solution–rationale structures, and the Discovery Engine [2] distils literature into structured knowledge artifacts and an integrated computational landscape. Both target the representation itself: what a source contains, recovered at scale. The decision studied here begins after that representation exists, and asks whether referent, construct, measurement, context, modality and attribution distinctions can govern a downstream mapping outcome and produce an obstruction when compatibility has not been earned. SciER [24] supplies scientific entity and relation annotations for datasets, methods and tasks in full text, which makes information extraction, discourse processing and entity/relation recovery an input to this work rather than part of it.

2.2 Schema-centric approaches

SciSchema.org [7] provides multidisciplinary, expert-reviewed scientific-process schemas with parameters, measurements, procedures and provenance, and is strong prior art for reusable scientific schema design. SCOPE and SCION [11] introduce an auditable benchmark and reference pipeline for evidence-linked schema induction and optional fusion from text, including conservative fusion behaviour. Executable Schema Contracts [13] discover an executable schema from heterogeneous sources and use it to drive provenance-aware construction and multi-source retrieval. Each of these produces a shared structure that a downstream consumer can rely on; none of them is asked whether a particular proposed mapping preserves the coordinates a specific scientific claim depends on. That is the discriminator evaluated here, over already-structured projections and including cases in which refusing a merge is the correct outcome.

2.3 Alignment, integration, and resolution

Ontology matching [8] studies correspondences such as equivalence, subsumption and disjointness between ontology entities, and language-model systems such as LLMATCH [23] add modern schema-preparation, candidate-selection and column-alignment machinery. Surveys of knowledge-graph construction [10] cover heterogeneous-source integration, ontology and schema matching, entity resolution, fusion, provenance and quality assurance, while reviews of scholarly knowledge graphs [19] cover extraction, construction, refinement and use of semantic scholarly infrastructures. The distinction pursued here is that a mapping decision is bound to several scientific coordinates at once, and that the exact source, provenance and non-preserved-coordinate information is retained with the resulting decision. Entity resolution and fusion are especially useful pressure cases, because a record-level match may be appropriate at exactly the moment a claim-level merge remains unsafe on construct, measurement, context or modality grounds.

2.4 Claim comparison and stance

Stance detection assigns target-relative labels such as support or opposition; surveys document its use in fact-checking and misinformation settings [9]. Apparent scientific disagreement, however, frequently arises from different context, modality, measurement or attribution coordinates rather than from a genuine opposition, and a mapping record that collapses every difference into one generic contradiction label discards the reason. Preserving that reason, rather than detecting stance, is what is evaluated here.

2.5 Scientific-variable semantics and interoperability

The I-ADOPT interoperability framework [16] provides a community-developed semantic decomposition for machine-readable observable and scientific-variable descriptions, and a recent benchmark builds an expert-annotated corpus of more than one hundred scientific variables and reports that large language models still struggle to construct those structured representations accurately [15]. FAIR 2.0 separates terminological, ontological, referential, schema and logical interoperability, and argues that semantic interoperability requires explicit mappings rather than one privileged ontology [21]. Recent provenance work goes further: DEC reads provenance-bearing statements as epistemic stance and preserves source-relative disagreement instead of collapsing every attributed assertion into a single factual core [20]. Structured variable semantics, generic semantic interoperability and provenance-aware pluralism are therefore supplied by this literature. The residual question is claim-relative and downstream of all of it: once a strong semantic integration stack has produced structured candidate mappings, what further conditions license an assertion of *scientific identity*? The decision contract studied here consumes ontology and schema alignment, provenance, plural-view machinery and consistency checks as inputs, and returns a glue, obstruction, plural-view or unresolved outcome from load-bearing construct, measurement, context, provenance-stance and global-composition conditions.

2.6 Measurement equivalence and construct validity

Measurement-equivalence and replication literatures established long ago that nominally identical outcome measures need not be comparable across settings, and that equivalence can require explicit testing [12]. The mechanism here makes measurement and operationalization one explicit coordinate of the mapping record and allows an unresolved measurement relation to block a global merge on its own.

2.7 Literature-based discovery and pluralism

Swanson’s literature-based discovery lineage [17, 18] shows how complementary literatures can expose implicit connections without direct citation links. The bounded requirement added here is on the bridge concept itself: it must retain compatible scientific meaning across the two source-local projections, or the bridge stays an obstruction rather than being promoted to a scientific integration. Scientific pluralism has similarly argued that multiple representations may remain legitimate without collapsing into one complete account [14, 5]; the computational question is whether an integration record can carry a plural or obstructed outcome, with exact source lineage, instead of forcing a canonical merged statement.

2.8 Scientific retrieval and cross-domain synthesis

OpenScholar [1] performs retrieval-augmented synthesis over the scientific literature at scale, and BioSage [22] combines retrieval, specialized agents, cross-disciplinary translation and reasoning for scientific tasks. Retrieval, synthesis, agent orchestration and cross-domain translation are baseline capability for the problem, as are language-model-assisted schema matching and induction. The experiment reported here deliberately holds all of that fixed and varies only the decision rule applied after scientific projections are already structured, so that the construct-validity question can be asked without a retrieval confound.

The empirical claim is correspondingly narrow. It is not that this combination solves raw-text scientific integration end to end; it is that on a disjoint, prospectively frozen public-reference holdout of already-structured scientific projections, the typed mapping calculus reduced false merges against flat predicate canonicalization while satisfying the false-split guard declared in advance. Raw-text extraction, full expert construct validity, recoverability and downstream utility are separate questions with separate evidence requirements.

3 Method

This section defines the scientific mapping objects the public-reference evaluation is defined over. The definitions are data-level: they do not depend on a particular language model or retrieval backend.

The implementation evaluated here is the ORION knowledge-integration layer; it is the artifact under test, and the rest of the paper describes the mapping rule rather than the product, so that each claim reads as a claim about the rule. Its source, protocols and evidence archives are listed in Section 10.

3.1 Three-valued outcomes

Every mapping decision measured in this paper is three-valued rather than binary. A proposed mapping may be determined admissible, determined inadmissible, or left *unresolved*; a coordinate relation may be determined compatible, determined incompatible, or left unresolved. An outcome that could not be determined is recorded and counted separately from an outcome determined to be false, and the two are never merged. The distinction is load-bearing rather than decorative: missing source authority, an absent coordinate or an unsupported relation produces an unresolved outcome, and an unresolved outcome is not evidence that the mapping is wrong. Throughout the paper, an outcome described as undetermined is this third value, carried in the evidence record and in every downstream count, never silently resolved to a negative.

3.2 Source and source-local projection

Let a source record s bind an external document or structured source item to an exact identity, version or revision, locator, and content commitment. The relevant scientific meaning is represented as a source-local projection

$$P_s = \langle R_s, C_s, M_s, X_s, \rho_s, \mu_s, a_s, \delta_s, \Pi_s \rangle, \quad (1)$$

with coordinates for referent R_s , construct C_s , measurement/operationalization M_s , context X_s , polarity ρ_s , modality μ_s , attribution a_s , discourse relation δ_s , and assumptions/unresolved information Π_s .

Each coordinate is a typed hypothesis supported by the source authority available for that case, not a value filled merely to complete a schema. In the public-reference datasets, unsupported coordinates remain absent or unresolved; the study does not ask a model to infer missing gold coordinates from raw papers.

3.3 Coordinate-wise comparison

A comparison between projections P_a and P_b is a vector of typed coordinate relations rather than a scalar similarity score:

$$q(P_a, P_b) = \langle \text{ref}(R_a, R_b), \text{const}(C_a, C_b), \text{meas}(M_a, M_b), \text{ctx}(X_a, X_b), \text{pol}(\rho_a, \rho_b), \text{mod}(\mu_a, \mu_b), \text{attr}(a_a, a_b), \text{disc}(\delta_a, \delta_b) \rangle. \quad (2)$$

Relations can express compatibility, incompatibility, conditional transformability, or unresolved evidence according to the frozen semantic taxonomies. A proposed relation does not acquire authority merely from lexical or embedding similarity.

3.4 Typed mapping and preservation conditions

A typed mapping m states which coordinate relations it licenses and under what preservation conditions. Its record contains:

- the proposed relation (for example identity, equivalence, conditional transform, subsumption, association-only, no-licensed-mapping, or unresolved);
- the coordinates the mapping claims to preserve;
- conditions that must hold for the mapping to be used; and
- coordinates or information that are not preserved and therefore may not be silently transferred into a global statement.

This representation is deliberately stricter than a successful match score: a mapping can be structurally plausible and still be inadmissible for the scientific claim if a load-bearing condition is absent.

3.5 Gluing, obstruction, and unresolved mapping

Two or more projections may be *glued* into one global statement only when the coordinate relations the mapping requires are compatible and its preservation conditions are satisfied. If a required relation is incompatible, the correct output is a typed *obstruction* naming the failing coordinate or condition. If the available evidence is insufficient, the relation stays unresolved rather than being forced to merge or split. Multiple locally valid but mutually incompatible views may remain a *plural view*.

The confirmatory experiment tests this decision semantics directly. Its strongest discrimination is on polarity, modality, attribution and context cases, where flat predicate canonicalization merges claims that the typed rule keeps separate. It does not identify the necessity of every coordinate; coordinate necessity requires targeted cases that the current datasets do not contain.

3.6 Provenance and recoverable records

Every case retains the external source identity and the frozen record used to derive the expected mapping decision. A mapping result retains the source records and the coordinate conditions supporting its terminal. That provenance is what makes the reported result independently replayable.

The wider design goes further and defines a global portrait

$$\mathcal{G} = \langle \{P_i\}, \mathcal{M}, \sigma \rangle, \quad (3)$$

where $\mathcal{M}$ stores accepted and rejected mappings and σ maps a global statement back to source projections and conditions. Whether a reader can in practice recover a generated portrait this way is a separate empirical question and is not answered by the current mapping study.

3.7 Search-universe feedback is outside the evaluated claim

The wider design also allows newly absorbed representations to change the search universe W and reopen future acquisition routes. That feedback is not manipulated or measured in the public-reference mapping experiment. No result in this paper is attributed to a “ W effect,” and no expansion of W is used as an ablation supporting the reported mapping result.

3.8 What is and is not claimed

Nothing here claims the first scientific knowledge graph, cross-domain retrieval-augmented system, schema matcher, provenance graph, entity resolver or literature-discovery method. The supported empirical object is narrower: over the frozen public-reference mapping cases, explicit scientific-coordinate conditions plus an obstruction state reduce false merges relative to flat predicate canonicalization while satisfying a conservative false-split guard.

Local exact worlds continue to falsify implementation semantics in controlled cases, in the manner established for source-local reconstruction and for typed contradiction detection [3, 4]. They do not establish external raw-text extraction quality, recoverability, downstream utility, or broad cross-domain generality; those remain separate prospective questions.

3.9 Method-structure projection

Scientific meaning coordinates describe what a source claims. Structural learning downstream of that also needs a source-local account of *how the source’s method transforms its problem*, so the representation is extended with a versioned method-structure projection. It is derived from a provenance-bound method-realization record of the kind produced by source-local epistemic reconstruction [3], and retains exact source, document, version and span identity, the target role and initial obstruction, assumptions and preconditions, representation changes and auxiliary objects where the source supports them, the transformation graph, invariants, progress and termination semantics, reconstruction to the original target, failure and negative-result statements, explicit author rationale when the source actually supplies it, and coordinates typed as unknown otherwise.

This projection stays source-local. In particular, the absence of an author rationale is not permission to invent one. The projection stores the exact realization and span digests that licensed each structured interpretation, and carries no authority to declare that two methods belong to the same formal family.

3.10 Cross-domain method alignment and plural views

Given two source-local projections, a claim-relative mapping μ may be proposed by comparing declared coordinates such as preconditions, assumptions, invariants, progress, termination and reconstruction. The alignment outcome is one of aligned, related, obstructed, or unresolved. The record preserves both projection identities, source-local names and meanings, unmatched assumptions, coordinate-level support and obstruction, and the exact reasons for refusing a merge. An aligned outcome means only that the declared source-local coordinates agree for the proposed mapping purpose; establishing formal equivalence is outside this mechanism’s authority.

A clean vocabulary-changing example is numerical bisection against monotone threshold calibration. Their operation names differ, but under the narrow purpose “bounded monotone localization” both require an ordered domain and a bracketed monotone target, preserve a target-containing bracket, make progress by shrinking the admissible interval, terminate at a tolerance, and reconstruct a representative in the original target space. A candidate alignment may therefore be recorded while each donor’s vocabulary and lineage is preserved.

A false-merge control uses the same-looking midpoint-and-discard-half surface sequence on a non-monotone response. The monotonicity and bracketing contract no longer holds and the invariant can fail, so the correct result is an obstruction; topology or lexical similarity cannot override that coordinate. A third control uses an opaque recovery procedure whose source does not establish progress or reconstruction semantics. Even when its known effects resemble a rollback method, the alignment stays unresolved rather than filling the missing coordinates from analogy.

Plural views are therefore first-class outcomes: structurally suggestive sources need not collapse into one global method. An obstruction certificate records exactly which coordinates prevent a merge and keeps both source projections recoverable.

3.11 Downstream structural-learning boundary

What this mechanism owns is source-local method projection and candidate plural alignment. Claim-relative structural reduction, substitution and method-family membership are outside it, as is any learned distribution over these versioned objects and any generated method structure proposed from such a distribution. Neither learned similarity nor generated structure changes a source projection or confers scientific authority on it. The bounded pilot reported in Section 5.7 tests only representation and alignment non-vacuity; it is not evidence of general extraction from arbitrary papers.

4 Public-reference mapping datasets

The empirical result in this paper uses two content-addressed *public-reference mapping datasets*. They contain already-structured scientific projections assembled from pinned external human and expert resources and from deterministic standards. They are not newly commissioned expert gold, and they do not evaluate raw-text extraction.

4.1 Initial public-reference set

The initial mapping study contains 32 cases generated from pinned revisions of MUSE, SciFact, and SciSchema. The portable gold archive is bound by content digest, and the evidence package binds the external-source revisions and evaluator identity while retaining a checksum manifest rather than redistributing restricted third-party text. This first result is treated as development evidence; it is not pooled into the confirmatory verdict.

Confirmatory case family	Cases
Different name / same referent	13
Polarity, modality, attribution, or context	13
Valid / invalid representation mapping	6
Total	32

Table 1: Case-family composition of the disjoint confirmatory public-reference holdout.

4.2 Disjoint confirmatory set

A second 32-case holdout was selected by an outcome-blind deterministic round-robin window over the sorted source pools, starting after the initial window and excluding all initial case identifiers. Its overlap with the first set is zero. Before any confirmatory system output was evaluated, the execution manifest froze the portable gold digest, the MUSE, SciFact and SciSchema revisions, evaluator component identities, bootstrap seed, margins, and terminal rule. Every digest named in that freeze is listed in Section 10.

Table 1 gives the composition. The 13 different-name/same-referent cases come from the MUSE cross-domain stratum. The 13 polarity/modality/attribution/context cases come from a scientific-claim-verification stratum. The six representation-mapping cases span materials, physics, and psychology and use deterministic standards where appropriate. Each record retains its source authority class and exact frozen input identity.

4.3 What the public-reference labels mean

The evaluated decision is a mapping-level terminal over already-structured projections. Depending on the source relation and the tested rule, the correct outcome can permit integration, forbid integration, or preserve an unresolved relation. The study evaluates false merges and false splits under these frozen cases. It does not claim that a model recovered the semantic coordinates from raw papers.

The confirmatory set is particularly discriminating for the family of polarity, modality, attribution and context cases: flat predicate canonicalization false-merges 6 of 13 cases (rate 0.4615), where coordinate-governed mapping makes no false merges. The different-name/same-referent and representation-mapping strata are clean controls on this holdout: all evaluated rules are correct there. That localization is reported rather than hidden, because the observed pooled effect is not evidence that every coordinate is necessary.

4.4 The broader expert-atlas design is not the dataset used here

A prospective expert-atlas design is preserved alongside this study: a protocol, annotation schema, handbook, adjudication policy, and 32 seed-to-gold templates spanning a broader eight-family design. Those templates are not expert gold. No independent dual annotation, per-coordinate inter-annotator agreement, specialist adjudication, or prospectively bound raw-text model run exists for that broader object, so its outcome remains undetermined and it is excluded from the claim made here rather than being presented retrospectively as the dataset used.

4.5 Data and licensing boundary

The public-reference package stores source identities, pinned revisions, locators, provenance and content commitments needed for replay. Third-party text is not vendored where redistribution is not permitted. A final submission archive will freeze the exact repository subject and stable public release URL for the legally shareable package; the absence of a permanent archive at the current drafting stage does not change the scientific result.

5 Evaluation

The empirical claim in this paper is governed by an execution-frozen protocol that was created after the initial 32-case result but before any confirmatory system output was examined. The predecessor result was allowed to motivate replication; it was not allowed to change the second holdout, the evaluator, the margins, or the pass criteria. The broader raw-text and expert-gold design remains a separate, unexecuted study. Both freezes are identified in Section 10.

5.1 Primary confirmatory hypothesis

The frozen replication hypothesis has two non-compensatory conditions:

False-merge superiority: the coordinate-governed-minus-flat false-merge point estimate must be at most -0.05 , and the upper bound of the paired 95% bootstrap interval must be below zero.

False-split non-inferiority: the coordinate-governed-minus-exact-conservative false-split point estimate must be at most $+0.03$, and the upper bound of its paired 95% bootstrap interval must be at most $+0.03$.

The holdout passes only if both conditions hold. It contains 32 cases, has zero overlap with the initial 32-case set, and was bound by exact content digest before any confirmatory outcome was accessed.

5.2 Compared mapping rules

The confirmatory primary comparison uses three deterministic rules over the same already-structured mapping cases:

1. **Typed mapping calculus:** evaluates the scientific coordinates the case requires and can retain an obstruction rather than force a merge.
2. **Flat predicate canonicalization:** treats compatible surface predicates as sufficient merge evidence without the coordinate distinctions. This is the predeclared false-merge comparator.
3. **Exact-coordinate conservative control:** requires exact coordinate compatibility and can abstain where it lacks such a match. This is the predeclared false-split and non-inferiority control.

These are mapping-layer comparators, not attempts to reproduce a full external knowledge-graph, retrieval-augmented or schema-fusion system. Stronger schema-induction and provenance-contract baselines would be useful post-outcome construct-validity pressure, but they cannot rewrite a confirmatory status that was frozen beforehand.

5.3 Secondary ablations

The frozen mapping evaluator also reports descriptive paired ablations. The covered attacks are:

- **force compatibility without obstruction:** removes the explicit non-merge terminal and chooses a compatible mapping even where the coordinates prohibit it;
- **remove modality/polarity/attribution/discourse:** removes the coordinate group carrying the main confirmatory discrimination;
- **remove referent, construct, measurement, and temporal context:** retained as zero-effect controls on this particular holdout.

The ablations are secondary and descriptive. The protocol explicitly forbids turning their intervals into new confirmatory significance claims, and forbids reading zero effect on an under-supported stratum as evidence that a coordinate is unnecessary.

5.4 Metrics and statistics

The confirmatory decision uses case-level false-merge and false-split rates. Wilson 95% intervals describe individual rates. Paired differences use a 10,000-resample percentile bootstrap with frozen seed 20260817. The mapping rules are deterministic, so the protocol uses one execution per case rather than artificial stochastic repeats. Results are reported pooled and descriptively by case family and source stratum. The initial 32 cases are never pooled into the confirmatory verdict.

The primary decision rule is therefore fixed entirely by the two paired comparisons above. No post-outcome margin change, exclusion, reweighting, or replacement baseline can change whether the confirmatory run passed.

5.5 Custody and fail-closed rules

Before any confirmatory output was produced, the execution manifest bound the portable confirmatory gold digest, zero overlap with the predecessor gold, the exact MUSE, SciFact and SciSchema source revisions, the evaluator component identities, the bootstrap seed and pass margins, and a rule that any digest mismatch, source drift, evaluator drift or case overlap leaves the result undetermined rather than silently re-evaluating a different study.

The gold is not an input to the evaluated mapping rule. The evaluator consumes the candidate mapping result and the protected expected decision only after the system output has been sealed.

5.6 Reproducibility

The result is reproduced by deterministic public-reference scripts over immutable archives: the confirmatory analysis, publication figures and tables, source registry, evaluator identities, and an independent replay receipt are all retained, and two make targets regenerate the headline analysis and the publication artifacts from committed inputs. Section 10 names each of them.

This reproducibility claim does not extend to the unexecuted raw-text expert atlas, to full external retrieval or schema baselines, to generated-portrait recoverability, or to downstream answer quality. Each of those remains undetermined until its own prospective run exists.

5.7 Bounded method-structure pilot

The public-reference scientific-meaning experiment is not an evaluation of method-structure extraction, and is not offered as one. The larger expert-annotated method-structure corpus originally proposed for this extension—50–100 papers and algorithms across at least five fields, with independent annotation—was not available for this submission track and is therefore outside the empirical claim.

Instead, a prospectively frozen non-vacuity pilot uses six authored exact-ground-truth method pairs across numerical methods, graph algorithms and transactional workflows. The panel contains three structurally aligned pairs, two explicit obstructions and one source-insufficient case whose correct outcome is unresolved. Eight controlled corruptions delete or change an invariant, a reconstruction map, a dependency order, a precondition, failure semantics, lineage, an authority boundary, or an explicitly unknown progress coordinate. The typed representation and alignment path reproduces all six frozen pair labels and detects all eight material corruptions. A deliberately simple transcript and action-overlap comparator resolves only two of the six pair labels.

That supports a narrow claim only: the typed source-local objects are discriminating on a small exact-ground-truth panel and can preserve an explicit unresolved state. It does not estimate annotation agreement on real papers, false merge or split rates over natural scientific prose, end-to-end extraction quality, or the performance of any learner built downstream of these objects. Those remain undetermined until independently annotated evidence exists, and the pilot narrows the method-projection claim rather than establishing a general cross-domain method-learning result.

6 Post-saturation validation: representational compatibility is not scientific identity

The public-reference result answers a question about flat canonicalization. A stronger comparator makes a wider question testable: if a semantic integration stack has already produced well-structured candidate mappings, does that structure by itself authorize an assertion of scientific identity? After the public-reference result was frozen, a separate study was opened to test that without rewriting the original 32-case evidence. Its protocol was frozen before protected outcomes, and it credits scientific-variable decomposition, ontology and schema alignment, provenance-aware integration, plural and inconsistency representation, and generic consistency checking as prior work. The only disputed interface is whether those representational outputs authorize scientific identity on their own.

The scientific-identity contract battery. The protected battery contains 400 exact already-structured contracts: five scientific domains, eight mapping archetypes and ten disjoint variants per cell. The correct terminal for each contract is to glue, to record an obstruction, to record a plural view, or to leave the mapping unresolved.

Strong products. The primary comparator is granted strong structured semantic information: construct, measurement, context, provenance and explicit missingness. It correctly handles direct construct, measurement and context mismatches and missing load-bearing coordinates. Its systematic differences are three: it treats provenance as trace rather than as scientific stance, it forces a canonical choice for legitimate one-to-many mappings, and it accepts pairwise-compatible mappings without a global scientific-consistency check. A second comparator performs canonical matching and represents the usual merge-on-high-local-match pressure. A third is an ideal fully

typed product that receives exactly the same coordinates and rules as the layer under test and is required to tie; it is an expressivity boundary rather than a weak baseline.

Protected result. The claim-relative compatibility layer reached exact scientific-integration decisions on 400/400 protected cases, against 250/400 for the strong semantic product and 50/400 for the canonical matcher. The paired difference against the strong product was +0.375 with a domain-stratified 95% bootstrap interval of [0.3275, 0.4225]; McNemar discordance was 150–0. The layer made no false merges, retained 1.0 clean-merge accuracy, and the paired effect was +0.375 in each of the five exact domains. An independent implementation reconstructed all four arms and the canonical protected-row digest exactly.

Boundary result. The ideal typed product also reached 400/400, with zero decision mismatches against the layer under test. The result therefore does not establish inherent expressivity or an advantage from centralization. What it supports is a bounded authority distinction: in the tested exact contracts, representational compatibility is weaker than scientific identity authority, and a claim-relative compatibility layer adds value over a strong semantic integration product by preserving obstruction, plurality and unresolved state where canonical identity is not licensed. The battery does not establish raw-text extraction, live ontology-engine superiority, expert agreement on new corpora, or open-ended cross-domain integration; each of those remains undetermined.

7 Results

The evidence layers are deliberately kept separate. The public-reference projection-to-mapping protocol has been executed twice: an initial frozen 32-case study, followed by a disjoint, prospectively execution-frozen 32-case confirmation. The broader raw-text expert atlas, generated-portrait recoverability, downstream utility and full coordinate necessity are separate questions whose outcomes remain undetermined; they are not treated as missing pieces of the mapping result reported here. All numbers below come from immutable archived artifacts, listed in Section 10.

7.1 Initial public-reference mapping result

The public-reference builder used pinned external source authorities from MUSE, SciFact, and SciSchema and did not fill unsupported semantic coordinates with model guesses. The first deterministic build produced 32 already-structured mapping cases across five recorded strata and three case families: different-name/same-referent, polarity/modality/attribution/context, and valid/invalid representation mappings. A second freeze pass reproduced the portable gold artifact byte for byte.

On those 32 cases, the coordinate-governed comparison rule classified all expected relations correctly (accuracy 1.000), with false-merge rate 0.000 and false-split rate 0.000. Flat predicate canonicalization had accuracy 0.875 and false-merge rate 0.125, with no false splits. The paired false-merge difference was -0.125 with a 95% paired percentile-bootstrap interval of $[-0.250, -0.03125]$ (10,000 resamples; seed 20260817).

An exact-coordinate conservative control also produced zero false merges but abstained on 0.125 of cases and achieved accuracy 0.875. The initial result therefore motivated a prospective replication without itself becoming the final confirmatory authority.

7.2 Prospectively frozen confirmatory replication

A second, disjoint 32-case holdout was selected with the same outcome-blind round-robin construction at selection offset 32, excluding every initial case identifier. It was frozen before any confirmatory output of any compared rule was examined; the execution manifest bound its portable digest, upstream source revisions, evaluator identities, bootstrap seed, margins, and the terminal rule before evaluation.

The rule declared in advance required both (i) a coordinate-governed-minus-flat false-merge difference at most -0.05 with the upper end of its paired 95% bootstrap interval below zero, and (ii) a coordinate-governed-minus-exact-conservative false-split difference no larger than $+0.03$ with the interval upper bound no larger than $+0.03$. The holdout met both conditions. Coordinate-governed mapping had false-merge rate 0.000 and false-split rate 0.000. Flat predicate canonicalization had false-merge rate 0.1875 and accuracy 0.8125. The paired false-merge difference was -0.1875 with 95% interval $[-0.34375, -0.0625]$. The false-split difference against the exact-coordinate control was 0.000 with interval $[0.000, 0.000]$. The confirmatory primary verdict was therefore a **pass** under the rule frozen before those outputs.

The effect is localized to the strata that discriminate the tested rules. Among the 13 polarity/modality/attribution/context cases, flat canonicalization false-merged 6/13 (0.4615), where coordinate-governed mapping false-merged none. The 13 different-name/same-referent cases and the six representation-mapping cases were correct under all compared rules. These strata are reported descriptively and are not pooled with the initial 32 cases to inflate the confirmatory sample size.

7.3 Covered ablations

On the confirmatory holdout, forcing compatibility without the explicit obstruction distinction increased false merges by $+0.1875$ (95% paired bootstrap interval $[+0.0625, +0.34375]$) and reduced accuracy by the same absolute amount. Removing the modality/polarity/attribution/discourse coordinate group produced the same false-merge increase on the covered claim-verification cases.

Removing referent, construct, measurement, or temporal context has zero measured effect on this holdout. Those zero effects are preserved as coverage limitations, not read as proof that the coordinates are unnecessary or dispensable. The confirmatory protocol treats all ablations as descriptive secondary analyses rather than as new significance claims.

7.4 Independent verification and local falsifiers

A separate implementation independently reproduced the confirmatory mapping result and its case-level decisions. The verification is bounded to the public-reference mapping claim; it does not promote the raw-text or downstream questions.

Synthetic exact-world tests separately exercise referent and construct identity, measurement equivalence, context, gluing, obstruction, plural views, preservation conditions, and typed mapping hypotheses. These remain engineering and semantic falsifiers rather than external empirical evidence.

7.5 What this result does not establish

The reported result does not establish:

- **Raw-text end-to-end superiority:** no prospectively bound raw-text extractor or model comparison against strong retrieval-augmented or schema systems supports this paper’s headline;
- **Generated-portrait recoverability or downstream utility:** no frozen reader-level or downstream task result is promoted here;
- **Universal coordinate necessity:** targeted cases for referent, construct, measurement, temporal context and additional obstruction families do not exist in these datasets;
- **An executed expert eight-family gold:** the broader seed templates and annotation protocol remain prospective rather than being relabelled as the public-reference dataset.

Each of those outcomes therefore remains undetermined, and stays visible as undetermined so that the replicated mapping result cannot be read as authority for the wider design. Any future positive has to arrive as a new prospectively frozen artifact rather than as a widening of the current claim in prose.

8 Limitations and ethics

8.1 Limitations

What is and is not owned here. Almost every mechanism this work depends on is prior art, and the boundary is worth stating once rather than repeatedly. Structured scientific-variable semantics, ontology and schema alignment, provenance modelling, provenance-aware pluralism, evidence-linked schema induction and fusion, entity resolution, stance and contradiction analysis, retrieval-augmented synthesis and generic consistency checking are all donor mechanisms, described and cited in Section 2. The residual is the claim-relative decision that runs after all of them: which coordinate differences license an assertion of scientific identity, and which oblige the record to keep an obstruction, a plural view or an unresolved relation visible instead. A result in favour of that decision layer is not a result in favour of the donor mechanisms, and none of the evidence here should be read as ownership of them.

Scope of the reported external result. The reported external result is the public-reference projection-to-mapping study, not an end-to-end expert atlas. It contains two separately frozen 32-case mapping sets assembled from pinned public human and expert resources; the confirmatory set is disjoint at the case level from the initial set. These cases are already-structured scientific projections. They therefore test mapping and obstruction semantics, not whether a model can reliably extract all required semantic coordinates from raw papers.

Scope of the post-saturation battery. The scientific-identity contract battery is a separate, prospectively frozen exact-contract study, not a pooled extension of the 32-case public-reference result. Its five domain labels denote heterogeneous scientific mapping contracts with mechanically defined ground truth; they are not executions of five live ontology or schema engines, and they do not substitute for newly commissioned expert agreement. The battery supports the bounded decision principle that representational compatibility can be weaker than scientific identity authority in the tested contracts. It does not establish ontology-engine-agnostic superiority or open-ended cross-domain integration.

Expert subjectivity and theory-ladenness. Scientific identity and measurement-equivalence judgments are theory-laden and may depend on disciplinary background. The prospective expert-atlas design mitigates this through a written handbook, independent double annotation, coordinate-wise agreement reporting, domain-expert escalation, and an explicit unresolved state. That expert-gold study has not been executed, so this manuscript reports no inter-annotator agreement and presents no seed template as expert gold.

Domain and case-family coverage. The public-reference route covers only the case families its pinned source authorities support. Its confirmatory effect is concentrated in polarity/modality/attribution/context cases; different-name/same-referent and representation-mapping cases are useful controls but do not discriminate the compared rules on that holdout. The successor battery broadens the exact contract families but still does not establish universal cross-domain adequacy or the necessity of every semantic coordinate. In particular, a coordinate’s usefulness in the exact successor does not imply that every scientific domain must encode that coordinate in the same representation.

Evolving terminology. Scientific terminology and operationalization change over time. The mapping study is bound to specific source revisions and does not establish robustness to temporal drift, historical conceptual change, or cross-lingual terminology.

Ontology and schema dependence. The semantic-coordinate representation is a design choice. Other representations may encode additional causal, probabilistic, temporal, or normative structure. The current evidence supports a bounded typed mapping and obstruction result and a bounded claim-relative compatibility successor; it does not establish that the chosen coordinate system is uniquely correct or optimal. An ideal product given the same scientific coordinates and rules ties the successor layer exactly.

Extraction error propagation. The text-to-scientific-meaning extraction layer remains unexecuted. Extraction errors could propagate into mapping, gluing and obstruction decisions. Because no prospectively bound raw-text extractor run supports the headline result, this paper reports no stage-attributed extraction error and claims no end-to-end superiority over raw-text, retrieval-augmented or schema systems.

Recoverability and downstream utility remain open. The design intends global statements to retain source projections, mapping paths, preservation conditions, and non-preserved coordinates. Local tests verify those mechanics, but the current external and exact-contract studies establish neither reader-level recoverability of generated portraits nor downstream scientific-answer improvement. Both remain undetermined pending their own frozen evaluations.

8.2 Ethics

The public-reference study reuses published research resources and collects no new human-participant data. The successor battery uses exact authored contracts rather than new human annotations. The broader expert-annotation protocol is prospective and has not been executed. An obstruction or a plural view should be presented as a provenance-bound record of incompatibility or of unresolved mapping, not as a judgment that one scientific perspective is intrinsically superior. Source lineage is retained so that a downstream reader can inspect the evidence and the conditions behind an integration decision.

9 Conclusion

This paper asked whether a global knowledge portrait — an integration of scientific claims across disciplines and measurement systems — can be built without treating surface or representational similarity as scientific identity. Three structural commitments answer that.

First, *projection before mapping*. Each source retains a local projection whose scientific meaning is represented by typed coordinates and bound to exact provenance before any cross-source decision is made. This keeps the question of what a source means separate from the question of which sources may be combined.

Second, *glue or obstruction*. Integration is decided over typed mapping hypotheses: projections are merged only when the relevant compatibility conditions hold; otherwise the record carries a typed obstruction, a contextual difference, a plural portrait, or an unresolved relation rather than silently flattening the distinction.

Third, *recoverability*. The intended global portrait exposes its construction — source projections, mapping path, preservation conditions, and non-preserved coordinates — so that a downstream consumer can audit a global statement back to its origins.

The formal projection and mapping semantics are verified by local exact worlds. The prospectively frozen public-reference route provides a bounded external check of the mapping layer without requiring a newly commissioned annotation team. On the disjoint 32-case confirmatory holdout assembled from pinned public human and expert resources, the typed comparison produced no false merges, where flat predicate canonicalization false-merged 0.1875 of cases; the paired difference was -0.1875 with a 95% bootstrap interval of $[-0.34375, -0.0625]$. Against the exact-coordinate conservative control, the false-split difference was 0.000 with interval $[0.000, 0.000]$, so the rule declared in advance passed.

The separate post-saturation battery tests a wider systems interpretation after granting the comparator strong semantic integration and current interoperability and provenance mechanisms. On 400 prospectively frozen exact already-structured contracts, the claim-relative compatibility layer made 400/400 exact scientific-integration decisions against 250/400 for a strong semantic product and 50/400 for canonical matching; the paired effect against the strong product was $+0.375$ with a domain-stratified 95% bootstrap interval of $[0.3275, 0.4225]$. It made no false merges and preserved 1.0 clean-merge coverage. An ideal product receiving exactly the same scientific coordinates and rules tied at 400/400, which is an explicit negative boundary against inherent-expressivity and centralization readings.

Together the two studies support a principle stronger than the original flat-canonicalization result: *representational compatibility is not itself scientific identity authority*. In the tested contracts, an explicit claim-relative compatibility layer can preserve obstruction, plurality and unresolved state after structured semantic alignment, rather than letting a technically valid mapping self-authorize a scientific merge.

That statement is deliberately narrower than end-to-end scientific integration. Variable semantics in the style of I-ADOPT, ontology, schema and entity alignment, provenance, provenance-aware pluralism and generic consistency checking all remain donor territory. Raw-text extraction, live ontology-engine generality, expert agreement on new domain corpora, recoverability of generated global portraits and downstream scientific utility remain undetermined until their own prospectively frozen evidence exists. Undetermined is the third outcome rather than a quiet negative: an outcome that could not be determined is reported separately from an outcome determined to be false, and no evidence gap in this paper is settled by reading one as the other.

10 Data and code availability

The implementation repository contains the mapping calculus, the frozen public-reference datasets and their builders, the deterministic evaluator, the successor battery, and the publication-layer regeneration commands. Every artifact named anywhere in this paper is listed here with its path, so that no repository location has to be carried in the narrative. Paths below are relative to the paper directory unless they begin with `src/` or `research/`.

10.1 Preregistration

The end-to-end study was preregistered and frozen before any outcome was accessed, under the identifier `P3.cross-domain-atlas.v1`. Its machine-readable freeze is `protocol/PROTOCOL_V1.json`, with the annotation schema, handbook and adjudication standard in `protocol/ANNOTATION_SCHEMA_V1.json`, `protocol/ANNOTATION_HANDBOOK_V1.md` and `protocol/ADJUDICATION_POLICY_V1.md`. That expert-gold design has not been executed; it is preserved so that a reader can see what was planned, and it supplies no evidence for anything reported here.

The executed mapping study has its own preregistration chain. The design freezes are `protocol/PUBLIC_REFERENCE_PROTocol_V1.json` and `protocol/PUBLIC_REFERENCE_PROTocol_V1_1.json`, and the execution freeze for the confirmatory holdout is `protocol/PUBLIC_REFERENCE_CONFIRMATORY_EXECUTION_V1.json`, registered under the identifier `P3.public-reference-mapping.v1.1-confirmatory`. A reader checking whether a reported analysis was planned or chosen after the fact should compare against that freeze; it is the reason every number in Section 7 can be traced to a decision made in advance of seeing it. The successor battery is frozen separately under `research/claim_expansion/p3/`, whose protocol, protected freeze, outcome-access receipt and result are `P3_X_PROTocol_V1.md`, `P3_X_PROTECTED_FREEZE_V1.json`, `P3_X_OUTCOME_ACCESS_RECEIPT_V1.md` and `P3_X_PROTECTED_RESULT_V1.json`.

10.2 Mechanism implementation

The coordinate comparison, bridge admissibility and mapping semantics defined in Section 3 are implemented in `src/orion/knowledge/semantics.py`; the public-reference builders, evaluator, analysis and publication projection are `src/orion/study/p3_public_reference_build.py`, `src/orion/study/p3_public_reference_holdout.py`, `src/orion/study/p3_public_reference.py`, `src/orion/study/p3_public_reference_analysis.py` and `src/orion/study/p3_public_reference_publication.py`. The independent evaluator used for replay is `src/orion/study/p3_independent_evaluator.py`.

10.3 Frozen public-reference datasets

The initial gold archive is `gold/adjudicated/public-reference-v1/PUBLIC_REFERENCE_GOLD_V1.jsonl` with its freeze manifest `gold/adjudicated/public-reference-v1/PUBLIC_REFERENCE_FREEZE_MANIFEST_V1.json`; its portable SHA-256 is `35f9e39b75ff53b7f0ec82cd03ebcaaa82509ee0aea3f5b96aac3fd62c854ed8`. The disjoint confirmatory holdout is the same pair of files under `gold/adjudicated/public-reference-v1.1-confirmatory/`, with portable SHA-256 `13a76c68c149c2552f3543babeca6e1ad5afe23c45ea9c0dc365c1445cf2782b`. The pinned external source authorities and their revisions are recorded in `gold/PUBLIC_REFERENCE_SOURCE_REGISTRY_V1.json`, and the sampling and authority rules are `gold/PUBLIC_REFERENCE_AUTHORITY_POLICY_V1.md` and `gold/GOLD_METHODODOLOGY_V1.md`. Restricted third-party text is not redistributed; the packages retain source identity, version, locator and content commitment instead.

10.4 Frozen results

The initial build report, evaluated summary, analysis, provenance record and checksum manifest are under `evidence/public-reference-v1/`; the external-source revisions and evaluator identity are bound in `evidence/public-reference-v1/PROVENANCE.env`. The confirmatory counterparts are under `evidence/public-reference-v1.1-confirmatory/`: the pre-outcome binding is `EXECUTION_MANIFEST.json`, the reported numbers are `CONFIRMATORY_ANALYSIS.json`, and the publication figures and tables are under `publication/` with their own `SHA256SUMS`. The bootstrap seed 20260817 and both pass margins are recorded in the execution manifest rather than chosen afterwards.

The independent reproduction of the confirmatory mapping result is `research/verification/records/P3.C5.confirmatory-mapping.json`; the successor battery’s independent verifier and its record are `research/claim_expansion/p3/verify_p3_x_independent.py` and `research/claim_expansion/p3/P3_X_INDEPENDENT_VERIFICATION_V1.json`. The bounded method-structure pilot of Section 5.7 has its protocol, runner and results in `research/extensions/p1-p3-structure/` and its receipt at `evidence/method-structure/BOUNDED_PILOT_RECEIPT_V1.md`.

The mapping from each result-bearing claim to the archived artifact that supports it is `CLAIM_LEDGER_V1.md`, with the manuscript map at `CLAIM_LEDGER_MANUSCRIPT_MAP_V1.md`.

10.5 Reproduction

Two make targets regenerate the reported analysis and the publication artifacts from committed inputs: `make paper03-public-reference` produces the headline analysis, and `make paper03-public-reference-publication` produces the publication figures and tables. Both are deterministic, and continuous integration re-runs the second on a clean checkout and requires the regenerated bytes to be identical to the committed ones. The short reproduction route is `REPRODUCE.md`.

Reproduction establishes only what the frozen artifacts establish. The unexecuted raw-text expert atlas, external retrieval and schema baselines, generated-portrait recoverability and downstream answer quality have no artifacts here, and their outcomes remain undetermined rather than absent by oversight.

References

- [1] Akari Asai et al. Openscholar: Synthesizing scientific literature with retrieval-augmented LMs. *arXiv preprint arXiv:2411.14199*, 2024.
- [2] Vladimir Baulin, Austin Cook, Daniel Friedman, Janna Lumiruuusu, Andrew Pashea, Shagor Rahman, and Benedikt Waldeck. The discovery engine: A framework for AI-driven synthesis and navigation of scientific knowledge landscapes. *arXiv preprint arXiv:2505.17500*, 2025.
- [3] Billy et al. ORION paper i: Recursive epistemic reconstruction via source-local projections, 2025. Project internal publication.
- [4] Billy et al. ORION paper ii: Interpretable scientific knowledge integration with typed contradiction detection, 2025. Project internal publication.
- [5] Nancy Cartwright. *The Dappled World: A Study of the Boundaries of Science*. Cambridge University Press, 1999.

- [6] Tsolia Cohen and Tom Hope. MUSE: A full-text cross-domain knowledge base of scientific problems, solutions, and rationales. *arXiv preprint arXiv:2608.10974*, 2026.
- [7] Jennifer D’Souza et al. Scischema.org: A multidisciplinary collection of schemas for structured scientific process descriptions. *arXiv preprint arXiv:2607.27955*, 2026.
- [8] Jérôme Euzenat and Pavel Shvaiko. *Ontology Matching*. Springer Berlin Heidelberg, 1 edition, 2007.
- [9] Momchil Hardalov, Arnav Arora, Preslav Nakov, and Isabelle Augenstein. A survey on stance detection for mis- and disinformation identification. In *Findings of the Association for Computational Linguistics: NAACL 2022*, pages 1259–1277. Association for Computational Linguistics, 2022.
- [10] Marvin Hofer, Daniel Obraczka, Alieh Saeedi, Hanna Köpcke, and Erhard Rahm. Construction of knowledge graphs: Current state and challenges. *Information*, 15(8):509, 2024.
- [11] Miaobo Hu, Xiaobo Guo, Shuhao Hu, Bokun Wang, Rui Chen, Xin Wang, Daren Zha, and Jun Xiao. SCOPE and SCION: A benchmark and an auditable reference pipeline for schema induction and fusion from text. *arXiv preprint arXiv:2607.21610*, 2026.
- [12] Sebastian Jilke, Nicolai Petrovsky, Bart Meuleman, and Oliver James. Measurement equivalence in replications of experiments: When and why it matters and guidance on how to determine equivalence. *Public Management Review*, 19(9):1293–1310, 2017.
- [13] Padmaja Jonnalagedda, Yuguang Yao, Xiang Gao, Hilaf Hasson, and Kamalika Das. Executable schema contracts: From automatic ingestion to multi-source retrieval. *arXiv preprint arXiv:2606.05415*, 2026.
- [14] Stephen H. Kellert, Helen E. Longino, and C. Kenneth Waters, editors. *Scientific Pluralism*, volume 19 of *Minnesota Studies in the Philosophy of Science*. University of Minnesota Press, 2006.
- [15] Barbara Magagna, Arvin Rastegar, Esteban González Guardia, Cristian Berrio, Stuart Chalk, José Manuel Gómez-Pérez, Christof Lorenz, Saurav Kumar, and Daniel Garijo. From scientific variables to knowledge graphs: The i-adopt benchmark. Semantic Web Journal submission 4005-5219, 2026. Major Revision as of July 2026.
- [16] Barbara Magagna, Ilaria Rosati, Maria Stoica, Sirko Schindler, Gwenaelle Moncoiffe, Anusuriya Devaraju, Johannes Peterseil, and Robert Huber. The i-adopt interoperability framework for fairer data descriptions of biodiversity. *arXiv preprint arXiv:2107.06547*, 2021.
- [17] Don R. Swanson. Fish oil, Raynaud’s syndrome, and undiscovered public knowledge. *Perspectives in Biology and Medicine*, 30(1):7–18, 1986.
- [18] Don R. Swanson and Neil R. Smalheiser. An interactive system for finding complementary literatures: A stimulus to scientific discovery. *Artificial Intelligence*, 91(2):183–203, 1997.
- [19] Shilpa Verma, Rajesh Bhatia, Sandeep Harit, and Sanjay Batish. Scholarly knowledge graphs through structuring scholarly communication: A review. *Complex & Intelligent Systems*, 9:1059–1095, 2023.

- [20] Fabio Vitali and Valentina Pasqual. Provenance-enhanced statements in knowledge graphs. *arXiv preprint arXiv:2606.15246*, 2026.
- [21] Lars Vogt, Philip Strömert, Nicolas Matentzoglou, Naouel Karam, Marcel Konrad, Manuel Prinz, and Roman Baum. Suggestions for extending the FAIR principles based on a linguistic perspective on semantic interoperability. *Scientific Data*, 12:688, 2025.
- [22] Svitlana Volkova, Peter Bautista, Avinash Hiriyanna, Gabriel Ganberg, Isabel Erickson, Zachary Klinefelter, Nick Abele, Hsien-Te Kao, and Grant Engberson. Cross-disciplinary knowledge retrieval and synthesis: A compound AI architecture for scientific discovery. *arXiv preprint arXiv:2511.18298*, 2025.
- [23] Sha Wang, Yuchen Li, Hanhua Xiao, Bing Tian Dai, Roy Ka-Wei Lee, Yanfei Dong, and Lambert Deng. LLMATCH: A unified schema matching framework with large language models. *arXiv preprint arXiv:2507.10897*, 2025.
- [24] Qi Zhang, Zhijia Chen, Huitong Pan, Cornelia Caragea, Longin Jan Latecki, and Eduard Dragut. SciER: An entity and relation extraction dataset for datasets, methods, and tasks in scientific documents. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 13083–13100, 2024.