

Adaptive State–Reasoning Co-Design under Matched Total Compute

ORION-P12

21 August 2026

Abstract

Test-time scaling is usually treated as a one-dimensional problem: decide how much more reasoning, sampling, search or verification to perform. State construction creates a second inference action. We formalize the resulting two-axis allocation problem under one budget and report a preregistered controlled run over 16 held-out generated families. The frozen joint arm scores 0.8582, versus 0.4631 and 0.4528 for the two named one-axis arms, and the run is byte-reproducible. A later comparator-capability audit changes the interpretation, not the historical bytes: the joint arm can emit four allocations while each one-axis arm can emit only two, and both losing-arm ceilings are below the winner’s achieved score. With identical four-action sets, the gain is +0.0408, the worst-family gain is +0.0020, and the original gate is not met. The comparison-validity adjudication therefore withholds P12A superiority authority. A prospectively frozen P12B then gives all arms the identical four actions and scores exact allocation. Across 32 independent family RNG blocks, the two-signal arm gains 0.253906 over the stronger one-signal arm (stratified family-block 95% bootstrap interval 0.251221 to 0.256653); every family and every fixed noise stratum passes. Current bounded authority is the controlled equal-action signal-complementarity result; naturalistic and external-system superiority remain open.

1 Introduction

Test-time computation has become an explicit design variable in modern reasoning systems. Systems allocate more tokens, samples, search nodes, verifier calls or iterative refinement to difficult instances. Recent work develops bandit, constrained-policy and learned adaptive allocation strategies, reinforcing a general lesson: uniform inference budgets waste computation when item difficulty is heterogeneous.

But difficulty itself has more than one source. Some tasks are difficult because the relevant structure is poorly exposed in the current representation. Others are difficult because substantial search or reasoning remains even after the right structure is visible. If a system spends all marginal budget on reasoning in an access-limited task, it reasons harder over the wrong state. If it spends all marginal budget on state construction in a reasoning-limited task, it repeatedly reorganizes information that was already accessible.

P12 asks:

Under one matched total budget, when should a system spend computation changing state, when should it spend computation reasoning over state, and can a prospective policy learn or exploit the difference?

The system is modeled as
raw/current state -> optional state construction -> downstream reasoning/search
-> verified outcome.

The paper makes four contributions.

1. **Two-axis inference formulation.** State construction and downstream reasoning are symmetric budgeted actions rather than free preprocessing plus paid reasoning.
2. **A strict comparator contract.** The joint policy must beat both adaptive-state-only and adaptive-reasoning-only policies at identical total resources, not merely a fixed baseline.
3. **A comparator-capability correction.** The protected run is reported with a later hostile adjudication showing that signal count and permitted allocations varied together.
4. **An equal-action successor.** P12B holds the four actions and budget fixed, varies only visible signals, and reports family-block uncertainty.

P12A’s superiority interpretation is withheld by its comparison-validity adjudication. P12B supplies the stronger controlled contract: equal total budget, equal action capability, and only then a variation in available signals. The current bounded authority is the equal-action signal-complementarity result.

2 Donor boundary and novelty

2.1 Adaptive test-time compute is prior-owned

Recent systems allocate inference compute dynamically based on predicted difficulty, value or resource constraints. Bandit formulations, constrained policy optimization, adaptive demonstration/generation strategies and “when to think” policies already own the primitive that different examples deserve different reasoning budgets. P12 therefore does not claim adaptive inference allocation itself.

2.2 Dynamic state construction is also prior-owned

Retrieval, compression, context selection, query-conditioned memory and structured-state construction already adapt what a model sees. P11 additionally supplies controlled evidence that construction can change accessibility. P12 does not claim dynamic state selection as a new primitive.

2.3 Residual after subtraction

The live residual is the **competition between those actions under one resource boundary**:

State construction and downstream reasoning are two places to spend test-time computation. A valid joint-allocation result must hold total resource fixed and strictly improve over policies allowed to adapt either axis alone.

In the current donor set, adaptive-compute methods optimize downstream reasoning/sampling/search or generation control; they do not make costed state construction and downstream reasoning symmetric decision variables under the same matched envelope and then require superiority over both one-axis adaptive controls.

3 Formal problem and why joint allocation can be strictly valuable

For item i , let:

- R_i be current/raw state;
- c_i be resource spent constructing/restructuring state;
- r_i be downstream reasoning/search resource;
- B_i be the total envelope;
- z_i be information available before the protected outcome;
- $Y_i(c_i, r_i)$ be verified success or quality.

The policy chooses (c_i, r_i) using z_i , subject to the common accounting contract. In the controlled benchmark the budget is scalar and exact: $c_i + r_i \leq B$. In a real system the resource is a vector and comparison is Pareto-based unless a cost scalarization is frozen before protected outcomes.

3.1 Policy classes

- `FIXED_STATE.FIXED_COMPUTE`: fixed allocation across items.
- `ADAPTIVE_STATE_ONLY`: may change state budget but not reasoning budget.
- `ADAPTIVE_REASON_ONLY`: may change reasoning budget but not state budget.
- `JOINT_STATE_REASONING`: may choose both under the same total envelope.
- `ORACLE_JOINT`: hindsight ceiling used only diagnostically.

Define

$$\text{joint_gain}(B) = Q_{\text{joint}}(B) - \max(Q_{\text{state_only}}(B), Q_{\text{reason_only}}(B)).$$

A positive P12 result requires $\text{joint_gain} > 0$ under the frozen comparison, not merely superiority to a fixed policy.

Consider a one-unit world containing two prospectively distinguishable regimes. In the access-limited regime success requires spending the unit on state construction; in the reasoning-limited regime success requires spending it on reasoning. A state-only adaptive policy cannot solve the latter and a reasoning-only adaptive policy cannot solve the former. A joint policy that sees the regime signal can spend the same unit at the valuable locus.

This existence argument is elementary and is not the empirical contribution. Its purpose is to identify the condition P12 must test: **heterogeneity in the location of marginal computation value**.

4 Protected matched-budget benchmark

4.1 Resource regimes

Each protected item has one hidden requirement:

- `EASY` = (0,0);
- `ACCESS` = (2,0);
- `REASON` = (0,2);
- `BOTH` = (1,1).

Every policy receives total budget $B=2$. Success is exact: allocated state and reasoning resources must meet both requirements. No arm receives extra budget and unused resource is not retrospectively reassigned.

4.2 Held-out families

The protected split contains **16 held-out families** $\times$ **512 items**. Family regime proportions vary, while a uniform mixture component prevents degenerate single-regime families. Signal noise σ_f ranges from 0.30 to 0.80 across families.

All adaptive arms receive the same pre-outcome signals

$$s_c = c_{req} + \text{Normal}(0, \sigma_f) \quad s_r = r_{req} + \text{Normal}(0, \sigma_f).$$

These signals contain no protected success outcome, verifier result or post-allocation feedback.

4.3 Frozen policies

- **FIXED_11**: always (1,1).
- **ADAPTIVE_STATE_ONLY**: choose (2,0) if $s_c \geq 1$, else (0,0).
- **ADAPTIVE_REASON_ONLY**: choose (0,2) if $s_r \geq 1$, else (0,0).
- **JOINT_FROZEN**: choose the feasible allocation in $\{(0,0), (1,1), (2,0), (0,2)\}$ nearest to (s_c, s_r) under squared Euclidean distance, with frozen tie order.
- **ORACLE_JOINT**: exact hindsight requirement, diagnostic only.

No policy is tuned on protected family outcomes.

5 Results

The historical runner printed its preregistered superiority terminal. The later comparison-validity adjudication withholds that claim authority. The current bounded authority comes from the prospectively frozen P12B equal-action result.

policy	mean verified success
JOINT_FROZEN	0.858154
FIXED_11	0.515503
ADAPTIVE_STATE_ONLY	0.463135
ADAPTIVE_REASON_ONLY	0.452759

The joint policy improves over the better one-axis adaptive policy by **mean** +0.334717, family-block 95% bootstrap CI [0.286008, 0.382693].

The **worst held-out family gain** is +0.158203. Joint versus fixed (1,1) gain is +0.342651 on average. Every allocation respects the two-unit budget, the oracle ceiling holds in every family, and two fresh executions produce the identical SHA-256

0194bc094f5696583533af5baae41e7c339902603d3706c8a1d2a78493f98947.

5.1 Why the historical contrast is not a signal-count result

Both one-axis policies use pre-outcome signals, but they may emit only two allocations while **JOINT_FROZEN** may emit four. **ADAPTIVE_STATE_ONLY** has a perfect-signal ceiling of 0.475464 and **ADAPTIVE_REASON_ONLY** a ceiling of 0.463623, both below the winner’s achieved 0.858154. The comparison therefore does not isolate the value of a second signal.

5.2 Regime interpretation

The large gain arises primarily because a one-axis policy’s restricted action set cannot serve the opposite-axis and jointly limited regimes at any signal value.

Giving one-signal policies the same four allocations yields mean gain 0.040771, family-block interval [0.031006, 0.050659], and worst-family gain 0.001953. The original positive gate is not met. This failure motivated a new frozen P12B rather than a post-hoc threshold change.

5.3 Prospectively frozen P12B

P12B gives all three arms the same four actions and budget two, and scores the exact required allocation. The independent unit is one family RNG block ($n=32$); 1,024 episodes within each block are technical observations. Mean two-signal gain over the stronger one-signal arm is 0.253906. The stratified family-block 95% bootstrap interval is [0.251221, 0.256653], minimum family gain is 0.196289, and every fixed sigma stratum exceeds 0.21. Two fresh subprocess payloads are byte-identical, supporting the controlled equal-action signal-complementarity claim.

6 Resource-accounting contract and statistical analysis

The scalar controlled budget is intentionally clean. Real systems require vector receipts including:

- compiler/retrieval/preprocessing model and operations;
- state tokens/bytes and memory traffic;
- downstream generated tokens/recurrent steps;
- search nodes and verifier calls;
- tool calls and external latency;
- cache/recovery cost;
- model identity/capacity;
- end-to-end latency and reproducible energy where available.

A joint policy cannot “win” by shortening the downstream trace while hiding expensive retrieval or compilation upstream. Likewise a reasoning-only baseline cannot receive a larger model or search cap. If resource vectors are incomparable, the result should be a quality–resource Pareto frontier rather than a post-hoc weighted score.

6.1 Statistical analysis

The protected unit is the held-out family RNG block. P12A has 16 family blocks. P12B has 32 independent family blocks, eight in each fixed noise stratum. Its primary deterministic 20,000-resample bootstrap samples families within each stratum, preserving the registered mixture. An unstratified 32-family bootstrap is reported only as sensitivity.

The analysis does not pool P12A’s 8,192 or P12B’s 32,768 technical episodes as independent domains. Hyperparameters are frozen before protected evaluation. The worst-family and minimum-stratum gains prevent a favorable mean from hiding a family- or stratum-level failure.

7 Relation to current adaptive-compute literature and limitations

Strategic test-time-compute allocation treats inference budget as a learnable or bandit decision across examples. Constrained policy approaches optimize accuracy under average compute. Adaptive in-context demonstration and generation methods jointly alter conditioning and generation effort. Recent “when to think” work likewise emphasizes selective reasoning to reduce unnecessary inference.

These results sharpen P12’s motivation: **adaptive inference is crowded; the novel discriminator must be where the resource can be spent**. P12A does not measure that discriminator because the action portfolio changes with signal count. P12B repairs the controlled equal-action contrast and supports bounded signal complementarity; it does not establish real-system resource-locus superiority.

The current bounded authority is the P12B equal-action result; the historical P12A comparator boundary remains separately content-bound.

7.1 Limitations and real-system promotion gate

1. The protected benchmark is a controlled resource world, not an LLM, prover or production agent.
2. The pre-outcome signals are constructed measurements of resource need. Real signal quality may be substantially worse.
3. Scalar units are commensurate by construction. Real compiler work, tokens, verifier calls and latency are heterogeneous.
4. The joint policy is a simple frozen nearest-allocation rule; the paper does not claim it is optimal.
5. P12B gives one- and two-signal policies an identical allocation set, but a real-system result must additionally include strong compute-only and state-only adaptive baselines.
6. A broad superiority claim requires at least one held-out real LLM/procedural domain or verifier-backed search domain under matched end-to-end resource receipts.
7. If real tasks overwhelmingly favor one resource locus, a simpler one-axis policy may be preferable; P12 predicts this as a regime condition rather than denying it.

8 Discussion and conclusion

P12 reframes test-time scaling as a **portfolio of computations**. “Think longer” is not the only adaptive action available to an intelligent system. It may be cheaper to parse, retrieve, compile, restructure or recover state so that less downstream search is required. Conversely, when state already exposes the relevant structure, additional preprocessing is wasteful and reasoning should receive the marginal budget.

P12A demonstrates the construction but not the key discriminator: equal total budget was real while equal action capability was not. P12B repairs that controlled estimand and finds a positive two-signal effect across the registered equal-action panel.

This leaves a concrete systems hypothesis for real agents: **test-time scaling curves may be two-dimensional, with state-work and reasoning-work measured on a common receipt and with action capability held fixed across signal ablations**.

Adaptive inference may need to decide not only **how much** computation to spend but **where** to spend it. P12 supplies the formulation, an exact failure analysis of its first empirical discriminator, and a positive prospectively frozen equal-action successor. Current authority is the bounded equal-action signal-complementarity result; the P12A comparison boundary remains historical, and matched real end-to-end validation is the next claim frontier.

9 References

- Zuo, B. & Zhu, Y. *Strategic Scaling of Test-Time Compute: A Bandit Learning Approach*. ICLR 2026.
- Zuo, B., Zhou, D. & Zhu, Y. *Adaptive Test-Time Compute Allocation with Evolving In-Context Demonstrations*. Findings of ACL 2026, 35156–35173. DOI: 10.18653/v1/2026.findings-acl.1754.
- Zhai, Z., Li, B., Xiao, B., Li, M. & Wang, X. *Adaptive Test-Time Compute Allocation for Reasoning LLMs via Constrained Policy Optimization*. arXiv:2604.14853, 2026.
- *Learning When to Think: Adaptive Reasoning for Test-Time Compute Allocation*. arXiv:2608.20256, 2026.
- P11 provides the controlled state-construction basis consumed by this paper; it does not transfer scientific authority to P12.